

Intro to Programming for Public Policy Week 8

Relational Databases and SQL

Eric Potash

May 21, 2018

Databases

Databases as Infrastructure

- ▶ Since they are not consumer-facing, the importance of databases as a technological infrastructure is not widely appreciated.
- ▶ They are a key component in most applications including websites, health records, supply chains, etc.

Why databases?

- ▶ Databases efficiently store and retrieve data.
 - ▶ Typical computers have about 8GB memory, a hard limit on what you can simultaneously load.
 - ▶ Databases allow efficient row-wise operations
 - ▶ Efficiently search hundreds of GB of data
 - ▶ Maintain one copy of the data and allow many people to access it (database server)

Example: American Time Use Survey (ATUS)

Respondents				
ID	Line	Children	Worked	...
1	1	0	1	...
2	1	1	0	...
3	1	2	1	...
4	1	0	1	...
5	1	3	1	...
6	1	0	0	...
7	1	1	1	...
...

Activities				
Resp.	Activity #	Code	Duration	...
1	1	010101	150	...
1	2	130124	45	...
1	3	010201	30	...
1	4	020201	10	...
1	5	110101	15	...
1	6	180501	20	...
1
2	1	010101	240	...
...

CPS				
Resp.	Line	Age	Education	...
1	1	38	38	...
1	2	35	0	...
1	3	4	1	...
2	1	53	1	...
2	2	58	0	...
...

* Primary Keys

Relational databases

Most databases use a *relational* model:

- ▶ Data are stored in *tables*
- ▶ Tables are linked using *keys*
- ▶ Interact with database using Structured Query Language (SQL)

Why *not* relational databases?

- ▶ Extremely large datasets (thousands of gigabytes)
- ▶ Not all data is relational (there are databases for that, too!)
- ▶ Not all tasks are database tasks (e.g. file manipulation, networking, image processing)

Normal Forms (E.F. Codd 1971)

1. Each record (row) should be 'atomic', i.e. non-divisible. A single record should not contain multiple divisible pieces.

Normal Forms (E.F. Codd 1971)

1. Each record (row) should be 'atomic', i.e. non-divisible. A single record should not contain multiple divisible pieces.
 - ▶ e.g. respondents table should not have a cell with the ages of all household members (instead there is a roster table)

Normal Forms (E.F. Codd 1971)

1. Each record (row) should be 'atomic', i.e. non-divisible. A single record should not contain multiple divisible pieces.
 - ▶ e.g. respondents table should not have a cell with the ages of all household members (instead there is a roster table)
2. No non-prime attribute of the table may be dependent on a proper subset of any candidate key.

Normal Forms (E.F. Codd 1971)

1. Each record (row) should be 'atomic', i.e. non-divisible. A single record should not contain multiple divisible pieces.
 - ▶ e.g. respondents table should not have a cell with the ages of all household members (instead there is a roster table)
2. No non-prime attribute of the table may be dependent on a proper subset of any candidate key.
 - ▶ e.g. the activity table, identified by the respondent and activity number, should not contain general information on the respondent.

Normal Forms (E.F. Codd 1971)

1. Each record (row) should be 'atomic', i.e. non-divisible. A single record should not contain multiple divisible pieces.
 - ▶ e.g. respondents table should not have a cell with the ages of all household members (instead there is a roster table)
2. No non-prime attribute of the table may be dependent on a proper subset of any candidate key.
 - ▶ e.g. the activity table, identified by the respondent and activity number, should not contain general information on the respondent.
3. The values in a row should refer uniquely to the key, not to a non-key attribute.

Normal Forms (E.F. Codd 1971)

1. Each record (row) should be 'atomic', i.e. non-divisible. A single record should not contain multiple divisible pieces.
 - ▶ e.g. `respondents` table should not have a cell with the ages of all household members (instead there is a `roster` table)
2. No non-prime attribute of the table may be dependent on a proper subset of any candidate key.
 - ▶ e.g. the `activity` table, identified by the respondent and activity number, should not contain general information on the respondent.
3. The values in a row should refer uniquely to the key, not to a non-key attribute.
 - ▶ e.g. the `activities` table should not include both the activity code and its definition.

Good practice

- ▶ Each table should contain a single logical element, without repetition.
- ▶ Meaningful and carefully selected *primary key*.
- ▶ Appropriate primary keys are what allow databases to work efficiently.

Database concepts

- ▶ Tables
- ▶ Variables
 - ▶ Types
- ▶ Constraints
 - ▶ Unique
 - ▶ Non-null
 - ▶ Primary key
 - ▶ Foreign key

The collection of tables, variables and constraints together define a database *schema*.

Structured Query Language (SQL)

SQL Overview

- ▶ The most popular model among databases is relational. These systems are called *RDBMSs*.
- ▶ SQL is the standardized language for working with an RDBMS.
- ▶ Originally intended to be user-facing and fairly “intuitive”.
- ▶ Despite standardization, the implementations (e.g. Oracle, PostgreSQL, SQLite) all have some differences.

Running SQLite

Download the [atus.sqlite](#) file then run this at the command line (in the same directory as the downloaded file):

```
$ sqlite3 atus.sqlite
SQLite version 3.13.0 2016-05-18 10:57:30
Enter ".help" for usage hints.
sqlite>
```

Navigating SQLite

The biggest differences between RDBMSs are accessing the metadata. To show tables in the database:

```
sqlite> .tables  
activities      cps              respondents    roster
```

Table schema

To show the schema of a table (its columns and types):

```
sqlite> .schema --indent cps
CREATE TABLE cps(
  case_id INTEGER,
  line_no INTEGER,
  family_income INTEGER,
  educational_attainment INTEGER,
  duration_of_masters INTEGER,
  years_in_college INTEGER,
  marital_status INTEGER,
  state_code INTEGER,
  age INTEGER,
  years_education INTEGER,
  PRIMARY KEY(case_id, line_no)
);
```

Possible types: integer, real (float), text (string)

Full schema

```
sqlite> .fullschema --indent
CREATE TABLE roster(
    case_id INTEGER,
    line_no INTEGER,
    how_related INTEGER,
    edited_age INTEGER,
    edited_sex INTEGER,
    PRIMARY KEY(case_id, line_no)
);
CREATE TABLE respondents(
    case_id INTEGER,
    line_no INTEGER,
    daily_time_alone INTEGER,
    daily_time_secondary_childcare_hh_or_own_children INTEGER,
    worked_last_week INTEGER,
    daily_time_with_family INTEGER,
    daily_time_with_friends INTEGER,
    daily_time_secondary_childcare_hh_children INTEGER,
    daily_time_with_friends INTEGER
```

Data dictionaries

- ▶ ATUS Data Dictionary:
 - ▶ family_income: p. 14
 - ▶ educational_attainment: p. 28
 - ▶ marital_status: p. 59
- ▶ ATUS Lexicon

Selecting columns

To see all of the data in the cps table:

```
SELECT * FROM cps;
```

To see just two of the columns:

```
SELECT marital_status, year_education FROM cps;
```

Notes:

- ▶ Queries end with a semi-colon
- ▶ Upper case keywords (SELECT, FROM, etc.) are a convention but not necessary.
- ▶ No whitespace requirements.

Query output

```
sqlite> SELECT * from cps;  
...  
20151212152498|1|11|43|-1|-1|1|39|37|17  
20151212152498|2|11|43|-1|-1|1|39|34|17  
20151212152498|3|11|-1|-1|-1|-1|39|8|-1  
20151212152498|4|11|-1|-1|-1|-1|39|5|-1
```

You may prefer the output after changing these settings:

```
sqlite> .mode columns  
sqlite> .headers on
```


LIMIT

Akin to head in the command line and pandas, use the SQL keyword LIMIT to restrict output:

```
SELECT * FROM cps LIMIT 10;
```

WHERE

To restrict to rows satisfying a certain condition use a WHERE clause.

- ▶ when years_education = -1 the education variable is missing:

```
SELECT years_education
FROM cps
WHERE years_education > -1;
```

- ▶ Use AND, OR, NOT to perform boolean logic, as in python:

```
SELECT years_education, state_code
FROM cps
WHERE years_education > -1 AND
      state_code = 17;           -- 17 is Illinois
```

- ▶ Comment to end of line using -- (akin to # in python)

ORDER BY

- ▶ Can sort results by a column (defaults to ascending):

```
SELECT * FROM cps  
ORDER BY family_income;
```

ORDER BY

- ▶ Can sort results by a column (defaults to ascending):

```
SELECT * FROM cps  
ORDER BY family_income;
```

- ▶ To make it descending:

```
SELECT * FROM cps  
ORDER BY family_income DESC;
```

ORDER BY

- ▶ Can sort results by a column (defaults to ascending):

```
SELECT * FROM cps  
ORDER BY family_income;
```

- ▶ To make it descending:

```
SELECT * FROM cps  
ORDER BY family_income DESC;
```

- ▶ Combine with LIMIT as to find top or bottom:

```
SELECT * FROM cps  
ORDER BY family_income DESC  
LIMIT 10;
```

COUNT

Very useful *aggregate* function is COUNT:

```
SELECT COUNT(*) FROM cps
WHERE educational_achievement > -1;
```

An aggregate function is one that returns a single row for the table.

COUNT details

- ▶ Could do `COUNT(family_income)` or `COUNT(age)` to count the number of non-null values in the given column.
- ▶ In this data nothing is null, missingness is represented by -1.
- ▶ `COUNT(*)` counts all rows.

More aggregates

- ▶ To calculate the average value of a column:

```
SELECT AVG(age) FROM cps;
```

- ▶ To get the maximum:

```
SELECT MAX(age) FROM cps;
```

- ▶ And minimum:

```
SELECT MIN(age) FROM cps;
```

- ▶ Combine them:

```
SELECT MIN(age), AVG(age), MAX(age)  
FROM cps;
```


Booleans and aggregation

To find the proportion of respondents whose educational attainment is available that have a bachelors degree:

```
SELECT AVG(educational_attainment > 42)
FROM cps
WHERE educational_attainment > 0;
```

GROUP BY

```
SELECT number_of_hh_children,  
       count(*),  
       AVG(daily_time_alone/60)  
FROM respondents  
GROUP BY number_of_hh_children;
```

number_of_hh_children	count(*)	AVG(daily_time_alone/60)
0	6360	7.10943396226415
1	1944	3.92181069958848
2	1734	3.42214532871972
3	599	2.779632721202
4	190	3.07368421052632
5	49	3.26530612244898
6	23	3.30434782608696
7	3	2.33333333333333
8	2	0.5
9	1	0.0

GROUP BY with ORDER BY

```
SELECT state_code,  
       COUNT(*) as Count,  
       AVG(educational_attainment > 42) AS Bachelors  
FROM cps  
WHERE educational_attainment > 0 -- i.e. defined  
GROUP BY state_code  
ORDER BY Bachelors DESC  
LIMIT 10;
```

Notes:

- ▶ In a GROUP BY, the WHERE clause filters the rows going in to the aggregation
- ▶ ORDER BY orders the resulting groups, not the rows in the source table (cps).
- ▶ AS gives the column an alias

HAVING

- ▶ In a GROUP BY, the WHERE clause filters the rows going in to the aggregation.
- ▶ To filter the resulting groups use HAVING:

```
SELECT state_code,  
       COUNT(*) as Count,  
       AVG(educational_attainment > 42) AS Bachelors  
FROM cps  
WHERE educational_attainment > 0 -- i.e. defined  
GROUP BY state_code  
HAVING Count > 100  
ORDER BY Bachelors DESC  
LIMIT 10;
```

JOIN

Relational databases are all about relationships! We can harness the relationships in the data with a JOIN query:

```
SELECT cps.case_id,  
       educational_attainment,  
       spouse_or_partner_present  
FROM respondents  
JOIN cps  
ON respondents.case_id = cps.case_id AND  
   cps.line_no = 1;
```

The result of this query has a row for each respondent (all respondents are in the CPS table).

JOIN with GROUP BY

The combination of JOIN and GROUP BY is very powerful:

```
SELECT
    educational_attainment,
    AVG(spouse_or_partner_present = 1) Married
FROM cps JOIN respondents
ON cps.case_id = respondents .case_id AND
    cps.line_no = 1
WHERE number_of_hh_children > 0
GROUP BY educational_attainment;
```