

Intro to Programming for Public Policy Week 9

More SQL and Python Integration

Eric Potash

May 24, 2018

More SQL

Saved scripts

Save a query to a file, e.g. p_bachelors.sql:

```
SELECT AVG(educational_attainment > 42)
FROM cps
WHERE educational_attainment > 0;
```

Then execute it from the command line like so:

```
$ sqlite3 atus.sqlite < p_bachelors.sql
0.242895945308274
```

GROUP BY

```
SELECT number_of_hh_children,  
       count(*),  
       AVG(daily_time_alone/60)  
FROM respondents  
GROUP BY number_of_hh_children;
```

number_of_hh_children	count(*)	AVG(daily_time_alone/60)
0	6360	7.10943396226415
1	1944	3.92181069958848
2	1734	3.42214532871972
3	599	2.779632721202
4	190	3.07368421052632
5	49	3.26530612244898
6	23	3.30434782608696
7	3	2.33333333333333
8	2	0.5
9	1	0.0

GROUP BY with ORDER BY

```
SELECT state_code,  
       COUNT(*) as Count,  
       AVG(educational_attainment > 42) AS Bachelors  
FROM cps  
WHERE educational_attainment > 0 -- i.e. defined  
GROUP BY state_code  
ORDER BY Bachelors DESC  
LIMIT 10;
```

Notes:

- ▶ In a GROUP BY, the WHERE clause filters the rows going in to the aggregation
- ▶ ORDER BY orders the resulting groups, not the rows in the source table (cps).
- ▶ AS gives the column an alias

HAVING

- ▶ In a GROUP BY, the WHERE clause filters the rows going in to the aggregation.
- ▶ To filter the resulting groups use HAVING:

```
SELECT state_code,  
       COUNT(*) as Count,  
       AVG(educational_attainment > 42) AS Bachelors  
FROM cps  
WHERE educational_attainment > 0 -- i.e. defined  
GROUP BY state_code  
HAVING Count > 100  
ORDER BY Bachelors DESC  
LIMIT 10;
```

JOIN

Relational databases are all about relationships! We can harness the relationships in the data with a JOIN query:

```
SELECT cps.case_id,  
       educational_attainment,  
       spouse_or_partner_present  
FROM respondents  
JOIN cps  
ON respondents.case_id = cps.case_id AND  
   cps.line_no = 1;
```

Use LEFT JOIN, RIGHT JOIN, FULL OUTER JOIN as needed.

The result of this query has a row for each respondent (all respondents are in the CPS table).

JOIN with GROUP BY

The combination of JOIN and GROUP BY is very powerful:

```
SELECT
    educational_attainment,
    AVG(spouse_or_partner_present = 1) Married
FROM cps JOIN respondents
ON cps.case_id = respondents.case_id AND
    cps.line_no = 1
WHERE number_of_hh_children > 0
GROUP BY educational_attainment;
```


Hours of sleep

```
SELECT case_id,  
       SUM( (activity_code = 010101) * duration) / 60  
       AS sleep  
FROM activities  
GROUP BY case_id
```

Subqueries

We can treat the result of one query as a temporary table in a new query:

```
SELECT case_id, edited_labor_force_status, sleep
FROM respondents JOIN
(SELECT case_id,
    SUM( (activity_code = 010101) * duration) / 60
    AS sleep
    FROM activities
    GROUP BY case_id
)
USING (case_id);
```

Adding a GROUP BY

```
SELECT edited_labor_force_status, AVG(sleep)
FROM respondents JOIN
(SELECT case_id,
      SUM( (activity_code = 010101) * duration) / 60
      AS sleep
FROM activities
GROUP BY case_id
)
USING (case_id)
GROUP BY 1;
```

edited_labor_force_status	AVG(sleep)
-----	-----
1	8.10439043512348
2	8.67241379310345
3	8.86080178173719
4	8.93861174267754
5	8.79434148088459

CREATE TABLE AS

```
CREATE TABLE sleep AS
    SELECT case_id,
           SUM( (activity_code = 010101) * duration) / 60
           AS sleep
    FROM activities
    GROUP BY case_id
;
```

```
SELECT edited_labor_force_status, AVG(sleep)
FROM sleep JOIN respondents
USING (case_id)
GROUP BY 1;
```

DROP TABLE

We can drop the sleep table:

```
DROP TABLE sleep;
```

Alternatively, you might find it useful to make the table temporary, so that it is not saved in the `atus.sqlite` database file:

```
CREATE TEMP TABLE sleep AS  
    ...  
;
```

CREATE TABLE

We can create a new database by passing a non-existent file to sqlite3:

```
$ sqlite3 new_db.sqlite
```

Now we can create a new table with the following SQL:

```
CREATE TABLE test (  
    id INTEGER PRIMARY KEY,  
    Name TEXT,  
    Birthday TEXT  
);
```

INSERT

We can insert a record into the table using an INSERT statement:

```
INSERT INTO test (Name, Birthday)
VALUES ('Eric', '1987-09-23');
```

We can check that it worked:

```
sqlite> SELECT * FROM test;
id      Name      Birthday
-----
1       Eric      1987-09-23
```

Because the id column is an INTEGER PRIMARY KEY, sqlite automatically populates it with an integer that it automatically increments.

INSERT multiple records

```
INSERT INTO test (Name, Birthday)
VALUES ('Washington', '1732-02-22'),
       ('Adams', '1735-10-30');
```

```
sqlite> SELECT * FROM test;
```

id	Name	Birthday
-----	-----	-----
1	Eric	1987-09-23
2	Washington	1732-02-22
3	Adams	1735-10-30

Import a CSV

```
CREATE TABLE chicago (  
    Last TEXT, First TEXT,  
    Position TEXT, Department TEXT, FP TEXT, SH TEXT,  
    Hours INTEGER, Salary REAL, Wage REAL  
);
```

Databases have special functions for loading data from various file formats:

```
.import salaries.csv chicago
```

Python Integration

sqlite3 module

sqlite3 module allows python to connect to sqlite and execute queries:

```
import sqlite3

con = sqlite3.connect('atus.sqlite')
result = con.execute("SELECT * FROM cps")
for row in result:
    print(row)
```

pandas and sqlite3

```
import sqlite3
import pandas as pd

con = sqlite3.connect('atus.sqlite')
cps = pd.read_sql('SELECT * FROM CPS',
                  con,
                  index_col='case_id')
```


Read query from file

Save marriage_by_education.sql:

```
SELECT
    educational_attainment,
    AVG(spouse_or_partner_present = 1) Married
FROM cps JOIN respondents
ON cps.case_id = respondents.case_id AND
    cps.line_no = 1
WHERE number_of_hh_children > 0
GROUP BY educational_attainment;
```

Plot query results

```
with open('marriage_by_education.sql') as f:  
    sql = f.read()  
m = pd.read_sql(sql, con,  
                index_col='educational_attainment')  
m.plot()
```



Child engagement example

- ▶ Want to explore the relationship between years of education and time spent with children.
 - ▶ Years education is in the cps table
 - ▶ Time spent with children is codes 301XX, 302XX, and 303XX.
- ▶ Also want to restrict to:
 - ▶ respondents with at least one child (`number_of_hh_children > 0`)
 - ▶ employed (`edited_labor_force_status < 3`)
 - ▶ These columns are in the respondents table.
- ▶ So we need to JOIN three tables: respondents, cps, and activities

Child engagement SQL

```
SELECT
    respondents.case_id,
    cps.years_education AS Education,
    SUM (( activity_code/100 IN (301, 302, 303))
          * duration /60.0) AS Engagement
FROM respondents
JOIN cps ON
    respondents.case_id = cps.case_id AND
    respondents.line_no = cps.line_no
JOIN activities ON
    respondents.case_id = activities.case_id
WHERE
    number_of_hh_children > 0 AND
    edited_labor_force_status < 3
GROUP BY respondents.case_id;
```


Child engagement Python

```
import sqlite3
import pandas as pd
from matplotlib import pyplot as plt

con = sqlite3.connect('atus.sqlite')
with open('child_engagement.sql') as f:
    sql = f.read()

df = pd.read_sql(sql , con)
ax = df.boxplot('Engagement', 'Education')
ax.set(ylim = (0, 7),
        xlabel="Parental Education (Years)",
        ylabel="Direct Engagement (Hours)")
```

Child engagement plot

